

dog_app

March 22, 2019

1 Convolutional Neural Networks

1.1 Project: Write an Algorithm for a Dog Identification App

In this notebook, some template code has already been provided for you, and you will need to implement additional functionality to successfully complete this project. You will not need to modify the included code beyond what is requested. Sections that begin with '**(IMPLEMENTATION)**' in the header indicate that the following block of code will require additional functionality which you must provide. Instructions will be provided for each section, and the specifics of the implementation are marked in the code block with a 'TODO' statement. Please be sure to read the instructions carefully!

Note: Once you have completed all of the code implementations, you need to finalize your work by exporting the Jupyter Notebook as an HTML document. Before exporting the notebook to html, all of the code cells need to have been run so that reviewers can see the final implementation and output. You can then export the notebook by using the menu above and navigating to **File -> Download as -> HTML (.html)**. Include the finished document along with this notebook as your submission.

In addition to implementing code, there will be questions that you must answer which relate to the project and your implementation. Each section where you will answer a question is preceded by a '**Question X**' header. Carefully read each question and provide thorough answers in the following text boxes that begin with '**Answer:**'. Your project submission will be evaluated based on your answers to each of the questions and the implementation you provide.

Note: Code and Markdown cells can be executed using the **Shift + Enter** keyboard shortcut. Markdown cells can be edited by double-clicking the cell to enter edit mode.

The rubric contains *optional* "Stand Out Suggestions" for enhancing the project beyond the minimum requirements. If you decide to pursue the "Stand Out Suggestions", you should include the code in this Jupyter notebook.

Step 0: Import Datasets

Make sure that you've downloaded the required human and dog datasets: * Download the [dog dataset](#). Unzip the folder and place it in this project's home directory, at the location /dog_images.

- Download the [human dataset](#). Unzip the folder and place it in the home directory, at location /lfw.

Note: If you are using a Windows machine, you are encouraged to use [7zip](#) to extract the folder.

In the code cell below, we save the file paths for both the human (LFW) dataset and dog dataset in the numpy arrays `human_files` and `dog_files`.

```
In [1]: import numpy as np
        from glob import glob

        # load filenames for human and dog images
        human_files = np.array(glob("/data/lfw/*/*"))
        dog_files = np.array(glob("/data/dog_images/*/*/*"))
        # print number of images in each dataset
        print('There are %d total human images.' % len(human_files))
        print('There are %d total dog images.' % len(dog_files))
```

There are 13233 total human images.

There are 8351 total dog images.

Step 1: Detect Humans

In this section, we use OpenCV's implementation of [Haar feature-based cascade classifiers](#) to detect human faces in images.

OpenCV provides many pre-trained face detectors, stored as XML files on [github](#). We have downloaded one of these detectors and stored it in the `haarcascades` directory. In the next code cell, we demonstrate how to use this detector to find human faces in a sample image.

```
In [2]: import cv2
        import matplotlib.pyplot as plt
        %matplotlib inline

        # extract pre-trained face detector
        face_cascade = cv2.CascadeClassifier('haarcascades/haarcascade_frontalface_alt.xml')

        # load color (BGR) image
        img = cv2.imread(human_files[0])

        # convert BGR image to grayscale
        gray = cv2.cvtColor(img, cv2.COLOR_BGR2GRAY)

        # find faces in image
        faces = face_cascade.detectMultiScale(gray)

        # print number of faces detected in the image
        print('Number of faces detected:', len(faces))

        # get bounding box for each detected face
        for (x,y,w,h) in faces:
            # add bounding box to color image
            cv2.rectangle(img, (x,y), (x+w,y+h), (255,0,0), 2)
```

```
# convert BGR image to RGB for plotting
cv_rgb = cv2.cvtColor(img, cv2.COLOR_BGR2RGB)

# display the image, along with bounding box
plt.imshow(cv_rgb)
plt.show()
```

Number of faces detected: 1



Before using any of the face detectors, it is standard procedure to convert the images to grayscale. The `detectMultiScale` function executes the classifier stored in `face_cascade` and takes the grayscale image as a parameter.

In the above code, `faces` is a numpy array of detected faces, where each row corresponds to a detected face. Each detected face is a 1D array with four entries that specifies the bounding box of the detected face. The first two entries in the array (extracted in the above code as `x` and `y`) specify the horizontal and vertical positions of the top left corner of the bounding box. The last two entries in the array (extracted here as `w` and `h`) specify the width and height of the box.

1.1.1 Write a Human Face Detector

We can use this procedure to write a function that returns `True` if a human face is detected in an image and `False` otherwise. This function, aptly named `face_detector`, takes a string-valued file path to an image as input and appears in the code block below.

```
In [3]: # returns "True" if face is detected in image stored at img_path
def face_detector(img_path):
```

```
img = cv2.imread(img_path)
gray = cv2.cvtColor(img, cv2.COLOR_BGR2GRAY)
faces = face_cascade.detectMultiScale(gray)
return len(faces) > 0
```

1.1.2 (IMPLEMENTATION) Assess the Human Face Detector

Question 1: Use the code cell below to test the performance of the `face_detector` function.

- What percentage of the first 100 images in `human_files` have a detected human face?
- What percentage of the first 100 images in `dog_files` have a detected human face?

Ideally, we would like 100% of human images with a detected face and 0% of dog images with a detected face. You will see that our algorithm falls short of this goal, but still gives acceptable performance. We extract the file paths for the first 100 images from each of the datasets and store them in the numpy arrays `human_files_short` and `dog_files_short`.

Answer:

```
In [4]: from tqdm import tqdm
```

```
human_files_short = human_files[:100]
dog_files_short = dog_files[:100]

#-#-# Do NOT modify the code above this line. #-#-#

## TODO: Test the performance of the face_detector algorithm
## on the images in human_files_short and dog_files_short.
all_face = np.vectorize(face_detector)
human_faces = all_face(human_files_short)
dog_faces = all_face(dog_files_short)
print('Percentage from human_files_short with detected human face: {:.2f}%'.format((sum(human_faces) / len(human_files_short)) * 100))
print('Percentage from dog_files_short with detected human face: {:.2f}%'.format((sum(dog_faces) / len(dog_files_short)) * 100))
```

Percentage from human_files_short with detected human face: 98.00%

Percentage from dog_files_short with detected human face: 17.00%

We suggest the face detector from OpenCV as a potential way to detect human images in your algorithm, but you are free to explore other approaches, especially approaches that make use of deep learning :). Please use the code cell below to design and test your own face detection algorithm. If you decide to pursue this *optional* task, report performance on `human_files_short` and `dog_files_short`.

```
In [ ]: ### (Optional)
        ### TODO: Test performance of another face detection algorithm.
        ### Feel free to use as many code cells as needed.
```

Step 2: Detect Dogs

In this section, we use a [pre-trained model](#) to detect dogs in images.

1.1.3 Obtain Pre-trained VGG-16 Model

The code cell below downloads the VGG-16 model, along with weights that have been trained on [ImageNet](#), a very large, very popular dataset used for image classification and other vision tasks. ImageNet contains over 10 million URLs, each linking to an image containing an object from one of 1000 categories.

```
In [5]: import torch
import torchvision.models as models

# define VGG16 model
VGG16 = models.vgg16(pretrained=True)

# check if CUDA is available
use_cuda = torch.cuda.is_available()

# move model to GPU if CUDA is available
if use_cuda:
    VGG16 = VGG16.cuda()
```

Downloading: "https://download.pytorch.org/models/vgg16-397923af.pth" to /root/.torch/models/vgg16-397923af.pth
100%|| 553433881/553433881 [00:06<00:00, 89882071.89it/s]

Given an image, this pre-trained VGG-16 model returns a prediction (derived from the 1000 possible categories in ImageNet) for the object that is contained in the image.

1.1.4 (IMPLEMENTATION) Making Predictions with a Pre-trained Model

In the next code cell, you will write a function that accepts a path to an image (such as 'dogImages/train/001.Affenpinscher/Affenpinscher_00001.jpg') as input and returns the index corresponding to the ImageNet class that is predicted by the pre-trained VGG-16 model. The output should always be an integer between 0 and 999, inclusive.

Before writing the function, make sure that you take the time to learn how to appropriately pre-process tensors for pre-trained models in the [PyTorch documentation](#).

```
In [6]: from PIL import Image
import torchvision.transforms as transforms

# Get image path
def image_path(img_path):
    image = Image.open(img_path).convert('RGB')
    trans = transforms.Compose([transforms.Resize(size=(244, 244)), transforms.ToTensor()])
    image = trans(image)[:3, :, :].unsqueeze(0)
    return image

def VGG16_predict(img_path):
    """
    Use pre-trained VGG-16 model to obtain index corresponding to
```

```

predicted ImageNet class for image at specified path

Args:
    img_path: path to an image

Returns:
    Index corresponding to VGG-16 model's prediction
'''

## Load and pre-process an image from the given img_path
## Return the *index* of the predicted class for that image

img = image_path(img_path)
if use_cuda:
    img = img.cuda()
ret = VGG16(img)

return torch.max(ret,1)[1].item() # predicted class index

```

1.1.5 (IMPLEMENTATION) Write a Dog Detector

While looking at the [dictionary](#), you will notice that the categories corresponding to dogs appear in an uninterrupted sequence and correspond to dictionary keys 151-268, inclusive, to include all categories from 'Chihuahua' to 'Mexican hairless'. Thus, in order to check to see if an image is predicted to contain a dog by the pre-trained VGG-16 model, we need only check if the pre-trained model predicts an index between 151 and 268 (inclusive).

Use these ideas to complete the dog_detector function below, which returns True if a dog is detected in an image (and False if not).

```

In [7]: ### returns "True" if a dog is detected in the image stored at img_path
def dog_detector(img_path):
    class_id = VGG16_predict(img_path)
    return (class_id in range(151,269))

```

1.1.6 (IMPLEMENTATION) Assess the Dog Detector

Question 2: Use the code cell below to test the performance of your dog_detector function.

- What percentage of the images in human_files_short have a detected dog?
- What percentage of the images in dog_files_short have a detected dog?

Answer:

```

In [9]: ### Test the performance of the dog_detector function
### on the images in human_files_short and dog_files_short.

def dog_counter(files):
    cnt = 0;
    for file in files:
        cnt += dog_detector(file)
    return cnt

```

```
print("Percentage from human_files_short with detected dog: {:.2f}%".format(100*dog_count/len(human_files_short)))
print("Percentage from dog_files_short with detected dog: {:.2f}%".format(100*dog_count/len(dog_files_short)))
```

```
Percentage from human_files_short with detected dog: 0.00%
Percentage from dog_files_short with detected dog: 100.00%
```

We suggest VGG-16 as a potential network to detect dog images in your algorithm, but you are free to explore other pre-trained networks (such as [Inception-v3](#), [ResNet-50](#), etc). Please use the code cell below to test other pre-trained PyTorch models. If you decide to pursue this *optional* task, report performance on `human_files_short` and `dog_files_short`.

```
In [ ]: ### (Optional)
        ### TODO: Report the performance of another pre-trained network.
        ### Feel free to use as many code cells as needed.
```

Step 3: Create a CNN to Classify Dog Breeds (from Scratch)

Now that we have functions for detecting humans and dogs in images, we need a way to predict breed from images. In this step, you will create a CNN that classifies dog breeds. You must create your CNN *from scratch* (so, you can't use transfer learning *yet!*), and you must attain a test accuracy of at least 10%. In Step 4 of this notebook, you will have the opportunity to use transfer learning to create a CNN that attains greatly improved accuracy.

We mention that the task of assigning breed to dogs from images is considered exceptionally challenging. To see why, consider that *even a human* would have trouble distinguishing between a Brittany and a Welsh Springer Spaniel.

Brittany	Welsh Springer Spaniel
----------	------------------------

It is not difficult to find other dog breed pairs with minimal inter-class variation (for instance, Curly-Coated Retrievers and American Water Spaniels).

Curly-Coated Retriever	American Water Spaniel
------------------------	------------------------

Likewise, recall that labradors come in yellow, chocolate, and black. Your vision-based algorithm will have to conquer this high intra-class variation to determine how to classify all of these different shades as the same breed.

Yellow Labrador	Chocolate Labrador
-----------------	--------------------

We also mention that random chance presents an exceptionally low bar: setting aside the fact

that the classes are slightly imbalanced, a random guess will provide a correct answer roughly 1 in 133 times, which corresponds to an accuracy of less than 1%.

Remember that the practice is far ahead of the theory in deep learning. Experiment with many different architectures, and trust your intuition. And, of course, have fun!

1.1.7 (IMPLEMENTATION) Specify Data Loaders for the Dog Dataset

Use the code cell below to write three separate [data loaders](#) for the training, validation, and test datasets of dog images (located at `dog_images/train`, `dog_images/valid`, and `dog_images/test`, respectively). You may find [this documentation on custom datasets](#) to be a useful resource. If you are interested in augmenting your training and/or validation data, check out the wide variety of [transforms](#)!

```
In [10]: import os
import numpy as np
import torch
import torchvision.transforms as transforms
from torchvision import datasets
from PIL import ImageFile
ImageFile.LOAD_TRUNCATED_IMAGES = True

### Data loaders for training, validation, and test sets

## Specification for appropriate transforms, and batch_sizes

batch_size = 20
num_workers = 0

data_dir = '/data/dog_images/'
train_dir = os.path.join(data_dir, 'train/')
valid_dir = os.path.join(data_dir, 'valid/')
test_dir = os.path.join(data_dir, 'test/')

# Normalization
standard_normalization = transforms.Normalize(mean=[0.485, 0.456, 0.406], std=[0.229, 0.229, 0.229])

# Transform - Random Resized Crop
data_transforms = {'train': transforms.Compose([transforms.RandomResizedCrop(224),
                                                transforms.RandomHorizontalFlip(),
                                                transforms.ToTensor(),
                                                standard_normalization]),
                   'val': transforms.Compose([transforms.Resize(256),
                                                transforms.CenterCrop(224),
                                                transforms.ToTensor(),
                                                standard_normalization]),
                   'test': transforms.Compose([transforms.Resize(size=(224,224)),
                                                transforms.ToTensor(),
                                                standard_normalization])}
```



```

    }

    # Data sets
    train_data = datasets.ImageFolder(train_dir, transform=data_transforms['train'])
    valid_data = datasets.ImageFolder(valid_dir, transform=data_transforms['val'])
    test_data = datasets.ImageFolder(test_dir, transform=data_transforms['test'])

    # Setup the loaders
    train_loader = torch.utils.data.DataLoader(train_data, batch_size=batch_size, num_workers=
    valid_loader = torch.utils.data.DataLoader(valid_data, batch_size=batch_size, num_workers=
    test_loader = torch.utils.data.DataLoader(test_data, batch_size=batch_size, num_workers=

    loaders_scratch = {'train': train_loader, 'valid': valid_loader, 'test': test_loader}

```

Question 3: Describe your chosen procedure for preprocessing the data. - How does your code resize the images (by cropping, stretching, etc)? What size did you pick for the input tensor, and why? - Did you decide to augment the dataset? If so, how (through translations, flips, rotations, etc)? If not, why not?

Answer:

- For resizing, RandomResizedCrop was used for training with also normalization on the images via the mean and standard deviation while transforms.Resize was used for testing. Also, Tensor size of 224x224 was applied on image size to allow proper RandomCrops of the original image.
- Dataset augment was decided using RandomResizedCrop and RandomHorizontalFlip to train more variations of the dataset and avoid overfitting.

1.1.8 (IMPLEMENTATION) Model Architecture

Create a CNN to classify dog breed. Use the template in the code cell below.

```

In [11]: import torch.nn as nn
import torch.nn.functional as F
import numpy as np
ImageFile.LOAD_TRUNCATED_IMAGES = True

# define the CNN architecture
num_classes = 133
class Net(nn.Module):
    ### TODO: choose an architecture, and complete the class
    def __init__(self):
        super(Net, self).__init__()
        ## Define layers of a CNN
        self.conv1 = nn.Conv2d(3, 32, 3, stride=2, padding=1)
        self.conv2 = nn.Conv2d(32, 64, 3, stride=2, padding=1)
        self.conv3 = nn.Conv2d(64, 128, 3, stride=1, padding=1)
        self.pool = nn.MaxPool2d(2, 2)
        self.fc1 = nn.Linear(7 * 7 * 128, 512)

```

```

        self.fc2 = nn.Linear(512, num_classes)
        self.dropout = nn.Dropout(p=0.2)

    def forward(self, x):
        x = self.pool(F.relu((self.conv1(x))))
        x = self.pool(F.relu((self.conv2(x))))
        x = self.pool(F.relu((self.conv3(x))))
        x = x.view(-1, 7*7*128)
        x = self.dropout(x)
        x = F.relu(self.fc1(x))
        x = self.dropout(x)
        x = self.fc2(x)
        return x

##-## You so NOT have to modify the code below this line. ##-##

# instantiate the CNN
model_scratch = Net()
print(model_scratch)
# move tensors to GPU if CUDA is available
if use_cuda:
    model_scratch.cuda()

Net(
  (conv1): Conv2d(3, 32, kernel_size=(3, 3), stride=(2, 2), padding=(1, 1))
  (conv2): Conv2d(32, 64, kernel_size=(3, 3), stride=(2, 2), padding=(1, 1))
  (conv3): Conv2d(64, 128, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
  (pool): MaxPool2d(kernel_size=2, stride=2, padding=0, dilation=1, ceil_mode=False)
  (fc1): Linear(in_features=6272, out_features=512, bias=True)
  (fc2): Linear(in_features=512, out_features=133, bias=True)
  (dropout): Dropout(p=0.2)
)

```

Question 4: Outline the steps you took to get to your final CNN architecture and your reasoning at each step.

Answer:

- Set Input Image Size to 224x224.
- For Convolution Layers, reduce the dimensions of (x, y) to (7x7) to inline with well established models.
- Depth was eventually increased to 128 to get as high as 10% accuracy.
- Layer 1 is Conv2d(3, 32, kernel_size=(3, 3), stride=(2, 2), padding=(1, 1)) with MaxPool: MaxPool2d(kernel_size=2, stride=2, padding=0, dilation=1, ceil_mode=False)
- Layer 2 is Conv2d(32, 64, kernel_size=(3, 3), stride=(2, 2), padding=(1, 1)) with MaxPool: MaxPool2d(kernel_size=2, stride=2, padding=0, dilation=1, ceil_mode=False)
- Layer 3 is Conv2d(64, 128, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1)) with MaxPool: MaxPool2d(kernel_size=2, stride=2, padding=0, dilation=1, ceil_mode=False)

- Layers are connected fully with ReLu activation and a dropout of 20%
- Layer 1 is (fc1): Linear(in_features=6272, out_features=512, bias=True) while Layer 2 is (fc2): Linear(in_features=512, out_features=133, bias=True)
- Loss is set to CrossEntropyLoss
- Optimizer is SGD (with learning rate = 0.05), and Output Size (num classes) is 133

1.1.9 (IMPLEMENTATION) Specify Loss Function and Optimizer

Use the next code cell to specify a [loss function](#) and [optimizer](#). Save the chosen loss function as `criterion_scratch`, and the optimizer as `optimizer_scratch` below.

```
In [12]: import torch.optim as optimization

        ### loss function
        criterion_scratch = nn.CrossEntropyLoss()

        ### select optimizer
        optimizer_scratch = optimization.SGD(model_scratch.parameters(), lr = 0.05)
```

1.1.10 (IMPLEMENTATION) Train and Validate the Model

Train and validate your model in the code cell below. [Save the final model parameters](#) at filepath `'model_scratch.pt'`.

```
In [15]: def train(n_epochs, loaders, model, optimizer, criterion, use_cuda, save_path):
        """returns trained model"""
        # initialize tracker for minimum validation loss
        valid_loss_min = np.Inf
        print_after = 100

        print("Training is now starting! Get some popcorn ;)\n")

        for epoch in range(1, n_epochs+1):
            # initialize variables to monitor training and validation loss
            train_loss = 0.0
            valid_loss = 0.0

            #####
            # train the model #
            #####
            model.train()

            for batch_idx, (data, target) in enumerate(loaders['train']):
                # move to GPU
                if use_cuda:
                    data, target = data.cuda(), target.cuda()
                ## find the loss and update the model parameters accordingly
                ## record the average training loss, using something like
                ## train_loss = train_loss + ((1 / (batch_idx + 1)) * (loss.data - train_loss))
```

```

        # initialize weights to zero
        optimizer.zero_grad()

        output = model(data)

        # calculate loss
        loss = criterion(output, target)

        # back prop
        loss.backward()

        # grad
        optimizer.step()

        train_loss += ((1 / (batch_idx + 1)) * (loss.data - train_loss))

        if batch_idx % print_after == 0:
            print(f'\tEpoch #{epoch}, Iteration #{batch_idx+1}, Loss: {train_loss}')

#####
# validate the model #
#####
model.eval()
for batch_idx, (data, target) in enumerate(loaders['valid']):

    # move to GPU
    if use_cuda:
        data, target = data.cuda(), target.cuda()

    output = model(data)
    loss = criterion(output, target)
    valid_loss += ((1 / (batch_idx + 1)) * (loss.data - valid_loss))

print('Epoch: {} \tTraining Loss: {:.6f} \tValidation Loss: {:.6f}'.format(epoch, train_loss, valid_loss))

## save the model if validation loss has decreased
if valid_loss < valid_loss_min:
    torch.save(model.state_dict(), save_path)
    print('Validation loss decreased ({:.6f} --> {:.6f}). Saving model ...'.format(valid_loss_min, valid_loss))
    valid_loss_min = valid_loss
    torch.save(model.state_dict(), save_path)
print("\n")

# return trained model
return model

```

```

# train the model
trained_epochs = 33
model_file = 'model_scratch.pt'

model_scratch = train(trained_epochs, loaders_scratch, model_scratch, optimizer_scratch,
                      criterion_scratch, use_cuda, model_file)

# load the model that got the best validation accuracy
model_scratch.load_state_dict(torch.load(model_file))

print("Training is now completed! Put the popcorn away ;)")

Training is now starting! Get some popcorn ;)

Epoch #1, Iteration #1, Loss: 4.49035120010376
Epoch #1, Iteration #101, Loss: 4.563215255737305
Epoch #1, Iteration #201, Loss: 4.570560455322266
Epoch #1, Iteration #301, Loss: 4.545999526977539
Epoch: 1          Training Loss: 4.545671          Validation Loss: 4.370954
Validation loss decreased (inf --> 4.370954).  Saving model ...

Epoch #2, Iteration #1, Loss: 4.632518768310547
Epoch #2, Iteration #101, Loss: 4.476144790649414
Epoch #2, Iteration #201, Loss: 4.474935054779053
Epoch #2, Iteration #301, Loss: 4.466214179992676
Epoch: 2          Training Loss: 4.463962          Validation Loss: 4.343661
Validation loss decreased (4.370954 --> 4.343661).  Saving model ...

Epoch #3, Iteration #1, Loss: 4.192032814025879
Epoch #3, Iteration #101, Loss: 4.409173488616943
Epoch #3, Iteration #201, Loss: 4.406735420227051
Epoch #3, Iteration #301, Loss: 4.409986972808838
Epoch: 3          Training Loss: 4.409357          Validation Loss: 4.202763
Validation loss decreased (4.343661 --> 4.202763).  Saving model ...

Epoch #4, Iteration #1, Loss: 4.2416486740112305
Epoch #4, Iteration #101, Loss: 4.352278232574463
Epoch #4, Iteration #201, Loss: 4.364720821380615
Epoch #4, Iteration #301, Loss: 4.366661548614502
Epoch: 4          Training Loss: 4.363083          Validation Loss: 4.118946
Validation loss decreased (4.202763 --> 4.118946).  Saving model ...

Epoch #5, Iteration #1, Loss: 3.754354953765869
Epoch #5, Iteration #101, Loss: 4.276883125305176

```

Epoch #5, Iteration #201, Loss: 4.277930736541748
Epoch #5, Iteration #301, Loss: 4.271509647369385
Epoch: 5 Training Loss: 4.271475 Validation Loss: 4.045333
Validation loss decreased (4.118946 --> 4.045333). Saving model ...

Epoch #6, Iteration #1, Loss: 4.438803672790527
Epoch #6, Iteration #101, Loss: 4.242791175842285
Epoch #6, Iteration #201, Loss: 4.214855670928955
Epoch #6, Iteration #301, Loss: 4.2111687660217285
Epoch: 6 Training Loss: 4.202993 Validation Loss: 4.026919
Validation loss decreased (4.045333 --> 4.026919). Saving model ...

Epoch #7, Iteration #1, Loss: 4.108998775482178
Epoch #7, Iteration #101, Loss: 4.160248279571533
Epoch #7, Iteration #201, Loss: 4.16623592376709
Epoch #7, Iteration #301, Loss: 4.152754306793213
Epoch: 7 Training Loss: 4.148472 Validation Loss: 3.982505
Validation loss decreased (4.026919 --> 3.982505). Saving model ...

Epoch #8, Iteration #1, Loss: 4.197998523712158
Epoch #8, Iteration #101, Loss: 4.085809230804443
Epoch #8, Iteration #201, Loss: 4.083142280578613
Epoch #8, Iteration #301, Loss: 4.096522808074951
Epoch: 8 Training Loss: 4.098329 Validation Loss: 3.911072
Validation loss decreased (3.982505 --> 3.911072). Saving model ...

Epoch #9, Iteration #1, Loss: 3.7527194023132324
Epoch #9, Iteration #101, Loss: 4.004271507263184
Epoch #9, Iteration #201, Loss: 4.042060852050781
Epoch #9, Iteration #301, Loss: 4.034053325653076
Epoch: 9 Training Loss: 4.027389 Validation Loss: 3.885652
Validation loss decreased (3.911072 --> 3.885652). Saving model ...

Epoch #10, Iteration #1, Loss: 3.7918620109558105
Epoch #10, Iteration #101, Loss: 3.9462153911590576
Epoch #10, Iteration #201, Loss: 3.9648027420043945
Epoch #10, Iteration #301, Loss: 3.980725049972534
Epoch: 10 Training Loss: 3.987282 Validation Loss: 4.015882

Epoch #11, Iteration #1, Loss: 3.7799296379089355
Epoch #11, Iteration #101, Loss: 3.9229090213775635
Epoch #11, Iteration #201, Loss: 3.90301775932312

Epoch #11, Iteration #301, Loss: 3.917524576187134
Epoch: 11 Training Loss: 3.919889 Validation Loss: 3.751864
Validation loss decreased (3.885652 --> 3.751864). Saving model ...

Epoch #12, Iteration #1, Loss: 3.5387673377990723
Epoch #12, Iteration #101, Loss: 3.8777763843536377
Epoch #12, Iteration #201, Loss: 3.868990182876587
Epoch #12, Iteration #301, Loss: 3.864940881729126
Epoch: 12 Training Loss: 3.869295 Validation Loss: 3.697682
Validation loss decreased (3.751864 --> 3.697682). Saving model ...

Epoch #13, Iteration #1, Loss: 3.017728567123413
Epoch #13, Iteration #101, Loss: 3.7979660034179688
Epoch #13, Iteration #201, Loss: 3.8244755268096924
Epoch #13, Iteration #301, Loss: 3.817145586013794
Epoch: 13 Training Loss: 3.815827 Validation Loss: 3.673929
Validation loss decreased (3.697682 --> 3.673929). Saving model ...

Epoch #14, Iteration #1, Loss: 3.642611265182495
Epoch #14, Iteration #101, Loss: 3.7233481407165527
Epoch #14, Iteration #201, Loss: 3.749748945236206
Epoch #14, Iteration #301, Loss: 3.7545359134674072
Epoch: 14 Training Loss: 3.759728 Validation Loss: 3.639833
Validation loss decreased (3.673929 --> 3.639833). Saving model ...

Epoch #15, Iteration #1, Loss: 3.071089506149292
Epoch #15, Iteration #101, Loss: 3.7255918979644775
Epoch #15, Iteration #201, Loss: 3.7178494930267334
Epoch #15, Iteration #301, Loss: 3.720290422439575
Epoch: 15 Training Loss: 3.722321 Validation Loss: 3.574137
Validation loss decreased (3.639833 --> 3.574137). Saving model ...

Epoch #16, Iteration #1, Loss: 3.86488676071167
Epoch #16, Iteration #101, Loss: 3.6759872436523438
Epoch #16, Iteration #201, Loss: 3.653087854385376
Epoch #16, Iteration #301, Loss: 3.677933931350708
Epoch: 16 Training Loss: 3.672444 Validation Loss: 3.624585

Epoch #17, Iteration #1, Loss: 3.286264419555664
Epoch #17, Iteration #101, Loss: 3.5914478302001953
Epoch #17, Iteration #201, Loss: 3.5900473594665527
Epoch #17, Iteration #301, Loss: 3.6135873794555664

Epoch: 17 Training Loss: 3.610248 Validation Loss: 3.500738
Validation loss decreased (3.574137 --> 3.500738). Saving model ...

Epoch #18, Iteration #1, Loss: 3.5148119926452637
Epoch #18, Iteration #101, Loss: 3.5487215518951416
Epoch #18, Iteration #201, Loss: 3.551593065261841
Epoch #18, Iteration #301, Loss: 3.5640525817871094

Epoch: 18 Training Loss: 3.558341 Validation Loss: 3.582408

Epoch #19, Iteration #1, Loss: 3.136556386947632
Epoch #19, Iteration #101, Loss: 3.4715702533721924
Epoch #19, Iteration #201, Loss: 3.5059027671813965
Epoch #19, Iteration #301, Loss: 3.503375291824341

Epoch: 19 Training Loss: 3.515627 Validation Loss: 3.479233
Validation loss decreased (3.500738 --> 3.479233). Saving model ...

Epoch #20, Iteration #1, Loss: 3.757024049758911
Epoch #20, Iteration #101, Loss: 3.4652233123779297
Epoch #20, Iteration #201, Loss: 3.4958155155181885
Epoch #20, Iteration #301, Loss: 3.4950413703918457

Epoch: 20 Training Loss: 3.499774 Validation Loss: 3.489099

Epoch #21, Iteration #1, Loss: 3.335376262664795
Epoch #21, Iteration #101, Loss: 3.450730800628662
Epoch #21, Iteration #201, Loss: 3.4476935863494873
Epoch #21, Iteration #301, Loss: 3.442173719406128

Epoch: 21 Training Loss: 3.449567 Validation Loss: 3.475384
Validation loss decreased (3.479233 --> 3.475384). Saving model ...

Epoch #22, Iteration #1, Loss: 3.5338141918182373
Epoch #22, Iteration #101, Loss: 3.324195146560669
Epoch #22, Iteration #201, Loss: 3.38456392288208
Epoch #22, Iteration #301, Loss: 3.3966166973114014

Epoch: 22 Training Loss: 3.407851 Validation Loss: 3.378596
Validation loss decreased (3.475384 --> 3.378596). Saving model ...

Epoch #23, Iteration #1, Loss: 3.4202163219451904
Epoch #23, Iteration #101, Loss: 3.316668748855591
Epoch #23, Iteration #201, Loss: 3.3444066047668457
Epoch #23, Iteration #301, Loss: 3.384953260421753

Epoch: 23 Training Loss: 3.392874 Validation Loss: 3.405930

Epoch #24, Iteration #1, Loss: 2.569550037384033
Epoch #24, Iteration #101, Loss: 3.273745536804199
Epoch #24, Iteration #201, Loss: 3.3321280479431152
Epoch #24, Iteration #301, Loss: 3.3390440940856934
Epoch: 24 Training Loss: 3.341266 Validation Loss: 3.423422

Epoch #25, Iteration #1, Loss: 3.320918560028076
Epoch #25, Iteration #101, Loss: 3.267322540283203
Epoch #25, Iteration #201, Loss: 3.3025524616241455
Epoch #25, Iteration #301, Loss: 3.299772262573242
Epoch: 25 Training Loss: 3.302087 Validation Loss: 3.496745

Epoch #26, Iteration #1, Loss: 2.8111181259155273
Epoch #26, Iteration #101, Loss: 3.241466760635376
Epoch #26, Iteration #201, Loss: 3.227729082107544
Epoch #26, Iteration #301, Loss: 3.2433764934539795
Epoch: 26 Training Loss: 3.259964 Validation Loss: 3.417350

Epoch #27, Iteration #1, Loss: 3.330902099609375
Epoch #27, Iteration #101, Loss: 3.1920969486236572
Epoch #27, Iteration #201, Loss: 3.1992762088775635
Epoch #27, Iteration #301, Loss: 3.205310821533203
Epoch: 27 Training Loss: 3.209116 Validation Loss: 3.480667

Epoch #28, Iteration #1, Loss: 2.362097978591919
Epoch #28, Iteration #101, Loss: 3.160122871398926
Epoch #28, Iteration #201, Loss: 3.1938650608062744
Epoch #28, Iteration #301, Loss: 3.206001043319702
Epoch: 28 Training Loss: 3.205211 Validation Loss: 3.284253
Validation loss decreased (3.378596 --> 3.284253). Saving model ...

Epoch #29, Iteration #1, Loss: 2.4436161518096924
Epoch #29, Iteration #101, Loss: 3.122122287750244
Epoch #29, Iteration #201, Loss: 3.157163619995117
Epoch #29, Iteration #301, Loss: 3.149078130722046
Epoch: 29 Training Loss: 3.149392 Validation Loss: 3.329807

Epoch #30, Iteration #1, Loss: 1.9952293634414673
Epoch #30, Iteration #101, Loss: 3.0769922733306885
Epoch #30, Iteration #201, Loss: 3.09073805809021
Epoch #30, Iteration #301, Loss: 3.1077113151550293

Epoch: 30 Training Loss: 3.116377 Validation Loss: 3.338471

Epoch #31, Iteration #1, Loss: 2.6330947875976562
Epoch #31, Iteration #101, Loss: 3.0712692737579346
Epoch #31, Iteration #201, Loss: 3.0889105796813965
Epoch #31, Iteration #301, Loss: 3.0796284675598145

Epoch: 31 Training Loss: 3.098001 Validation Loss: 3.294354

Epoch #32, Iteration #1, Loss: 3.588583469390869
Epoch #32, Iteration #101, Loss: 3.0308022499084473
Epoch #32, Iteration #201, Loss: 3.053354024887085
Epoch #32, Iteration #301, Loss: 3.0491809844970703

Epoch: 32 Training Loss: 3.050368 Validation Loss: 3.410954

Epoch #33, Iteration #1, Loss: 2.7967026233673096
Epoch #33, Iteration #101, Loss: 2.9777023792266846
Epoch #33, Iteration #201, Loss: 2.982584238052368
Epoch #33, Iteration #301, Loss: 3.013981819152832

Epoch: 33 Training Loss: 3.019012 Validation Loss: 3.281077

Validation loss decreased (3.284253 --> 3.281077). Saving model ...

Training is now completed! Put the popcorn away ;)

1.1.11 (IMPLEMENTATION) Test the Model

Try out your model on the test dataset of dog images. Use the code cell below to calculate and print the test loss and accuracy. Ensure that your test accuracy is greater than 10%.

```
In [16]: def test(loaders, model, criterion, use_cuda):

    # monitor test loss and accuracy
    test_loss = 0.
    correct = 0.
    total = 0.

    model.eval()
    for batch_idx, (data, target) in enumerate(loaders['test']):
        # move to GPU
        if use_cuda:
            data, target = data.cuda(), target.cuda()
        # forward pass: compute predicted outputs by passing inputs to the model
        output = model(data)
        # calculate the loss
```

```

        loss = criterion(output, target)
        # update average test loss
        test_loss += ((1 / (batch_idx + 1)) * (loss.data - test_loss))
        # convert output probabilities to predicted class
        pred = output.data.max(1, keepdim=True)[1]
        # compare predictions to true label
        correct += np.sum(np.squeeze(pred.eq(target.data.view_as(pred))).cpu().numpy())
        total += data.size(0)

    print('Test Loss: {:.6f}\n'.format(test_loss))

    print('\nTest Accuracy: %2d%% (%2d/%2d)' % (100. * correct / total, correct, total))

    # call test function
    test(loaders_scratch, model_scratch, criterion_scratch, use_cuda)

```

Test Loss: 3.489491

Test Accuracy: 18% (154/836)

Step 4: Create a CNN to Classify Dog Breeds (using Transfer Learning)

You will now use transfer learning to create a CNN that can identify dog breed from images. Your CNN must attain at least 60% accuracy on the test set.

1.1.12 (IMPLEMENTATION) Specify Data Loaders for the Dog Dataset

Use the code cell below to write three separate [data loaders](#) for the training, validation, and test datasets of dog images (located at dogImages/train, dogImages/valid, and dogImages/test, respectively).

If you like, **you are welcome to use the same data loaders from the previous step**, when you created a CNN from scratch.

```

In [17]: # Specify data loaders
        loaders_transfer = loaders_scratch.copy()

```

1.1.13 (IMPLEMENTATION) Model Architecture

Use transfer learning to create a CNN to classify dog breed. Use the code cell below, and save your initialized model as the variable model_transfer.

```

In [18]: import torchvision.models as models
        import torch.nn as nn
        #import torch.optim as optim

```

```

# Specify model architecture
model_transfer = models.resnet50(pretrained=True)

for p in model_transfer.parameters():
    p.requires_grad = False

model_transfer.fc = nn.Linear(2048, 133, bias=True)
fc_parameters = model_transfer.fc.parameters()

for p in fc_parameters:
    p.requires_grad = True

print (model_transfer)

if use_cuda:
    model_transfer = model_transfer.cuda()

```

Downloading: "https://download.pytorch.org/models/resnet50-19c8e357.pth" to /root/.torch/models/100%|| 102502400/102502400 [00:01<00:00, 78738428.06it/s]

```

ResNet(
  (conv1): Conv2d(3, 64, kernel_size=(7, 7), stride=(2, 2), padding=(3, 3), bias=False)
  (bn1): BatchNorm2d(64, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
  (relu): ReLU(inplace)
  (maxpool): MaxPool2d(kernel_size=3, stride=2, padding=1, dilation=1, ceil_mode=False)
  (layer1): Sequential(
    (0): Bottleneck(
      (conv1): Conv2d(64, 64, kernel_size=(1, 1), stride=(1, 1), bias=False)
      (bn1): BatchNorm2d(64, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
      (conv2): Conv2d(64, 64, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1), bias=False)
      (bn2): BatchNorm2d(64, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
      (conv3): Conv2d(64, 256, kernel_size=(1, 1), stride=(1, 1), bias=False)
      (bn3): BatchNorm2d(256, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
      (relu): ReLU(inplace)
      (downsample): Sequential(
        (0): Conv2d(64, 256, kernel_size=(1, 1), stride=(1, 1), bias=False)
        (1): BatchNorm2d(256, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
      )
    )
  )
  (1): Bottleneck(
    (conv1): Conv2d(256, 64, kernel_size=(1, 1), stride=(1, 1), bias=False)
    (bn1): BatchNorm2d(64, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
    (conv2): Conv2d(64, 64, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1), bias=False)
    (bn2): BatchNorm2d(64, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
    (conv3): Conv2d(64, 256, kernel_size=(1, 1), stride=(1, 1), bias=False)
    (bn3): BatchNorm2d(256, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
    (relu): ReLU(inplace)
  )
)

```

```

)
(2): Bottleneck(
  (conv1): Conv2d(256, 64, kernel_size=(1, 1), stride=(1, 1), bias=False)
  (bn1): BatchNorm2d(64, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
  (conv2): Conv2d(64, 64, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1), bias=False)
  (bn2): BatchNorm2d(64, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
  (conv3): Conv2d(64, 256, kernel_size=(1, 1), stride=(1, 1), bias=False)
  (bn3): BatchNorm2d(256, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
  (relu): ReLU(inplace)
)
)
(layer2): Sequential(
  (0): Bottleneck(
    (conv1): Conv2d(256, 128, kernel_size=(1, 1), stride=(1, 1), bias=False)
    (bn1): BatchNorm2d(128, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
    (conv2): Conv2d(128, 128, kernel_size=(3, 3), stride=(2, 2), padding=(1, 1), bias=False)
    (bn2): BatchNorm2d(128, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
    (conv3): Conv2d(128, 512, kernel_size=(1, 1), stride=(1, 1), bias=False)
    (bn3): BatchNorm2d(512, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
    (relu): ReLU(inplace)
    (downsample): Sequential(
      (0): Conv2d(256, 512, kernel_size=(1, 1), stride=(2, 2), bias=False)
      (1): BatchNorm2d(512, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
    )
  )
  (1): Bottleneck(
    (conv1): Conv2d(512, 128, kernel_size=(1, 1), stride=(1, 1), bias=False)
    (bn1): BatchNorm2d(128, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
    (conv2): Conv2d(128, 128, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1), bias=False)
    (bn2): BatchNorm2d(128, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
    (conv3): Conv2d(128, 512, kernel_size=(1, 1), stride=(1, 1), bias=False)
    (bn3): BatchNorm2d(512, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
    (relu): ReLU(inplace)
  )
  (2): Bottleneck(
    (conv1): Conv2d(512, 128, kernel_size=(1, 1), stride=(1, 1), bias=False)
    (bn1): BatchNorm2d(128, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
    (conv2): Conv2d(128, 128, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1), bias=False)
    (bn2): BatchNorm2d(128, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
    (conv3): Conv2d(128, 512, kernel_size=(1, 1), stride=(1, 1), bias=False)
    (bn3): BatchNorm2d(512, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
    (relu): ReLU(inplace)
  )
  (3): Bottleneck(
    (conv1): Conv2d(512, 128, kernel_size=(1, 1), stride=(1, 1), bias=False)
    (bn1): BatchNorm2d(128, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
    (conv2): Conv2d(128, 128, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1), bias=False)
    (bn2): BatchNorm2d(128, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)

```

```

        (conv3): Conv2d(128, 512, kernel_size=(1, 1), stride=(1, 1), bias=False)
        (bn3): BatchNorm2d(512, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
        (relu): ReLU(inplace)
    )
)
(layer3): Sequential(
  (0): Bottleneck(
    (conv1): Conv2d(512, 256, kernel_size=(1, 1), stride=(1, 1), bias=False)
    (bn1): BatchNorm2d(256, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
    (conv2): Conv2d(256, 256, kernel_size=(3, 3), stride=(2, 2), padding=(1, 1), bias=False)
    (bn2): BatchNorm2d(256, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
    (conv3): Conv2d(256, 1024, kernel_size=(1, 1), stride=(1, 1), bias=False)
    (bn3): BatchNorm2d(1024, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
    (relu): ReLU(inplace)
    (downsample): Sequential(
      (0): Conv2d(512, 1024, kernel_size=(1, 1), stride=(2, 2), bias=False)
      (1): BatchNorm2d(1024, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
    )
  )
)
  (1): Bottleneck(
    (conv1): Conv2d(1024, 256, kernel_size=(1, 1), stride=(1, 1), bias=False)
    (bn1): BatchNorm2d(256, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
    (conv2): Conv2d(256, 256, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1), bias=False)
    (bn2): BatchNorm2d(256, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
    (conv3): Conv2d(256, 1024, kernel_size=(1, 1), stride=(1, 1), bias=False)
    (bn3): BatchNorm2d(1024, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
    (relu): ReLU(inplace)
  )
)
  (2): Bottleneck(
    (conv1): Conv2d(1024, 256, kernel_size=(1, 1), stride=(1, 1), bias=False)
    (bn1): BatchNorm2d(256, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
    (conv2): Conv2d(256, 256, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1), bias=False)
    (bn2): BatchNorm2d(256, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
    (conv3): Conv2d(256, 1024, kernel_size=(1, 1), stride=(1, 1), bias=False)
    (bn3): BatchNorm2d(1024, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
    (relu): ReLU(inplace)
  )
)
  (3): Bottleneck(
    (conv1): Conv2d(1024, 256, kernel_size=(1, 1), stride=(1, 1), bias=False)
    (bn1): BatchNorm2d(256, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
    (conv2): Conv2d(256, 256, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1), bias=False)
    (bn2): BatchNorm2d(256, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
    (conv3): Conv2d(256, 1024, kernel_size=(1, 1), stride=(1, 1), bias=False)
    (bn3): BatchNorm2d(1024, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
    (relu): ReLU(inplace)
  )
)
  (4): Bottleneck(
    (conv1): Conv2d(1024, 256, kernel_size=(1, 1), stride=(1, 1), bias=False)

```

```

        (bn1): BatchNorm2d(256, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
        (conv2): Conv2d(256, 256, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1), bias=False)
        (bn2): BatchNorm2d(256, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
        (conv3): Conv2d(256, 1024, kernel_size=(1, 1), stride=(1, 1), bias=False)
        (bn3): BatchNorm2d(1024, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
        (relu): ReLU(inplace)
    )
(5): Bottleneck(
  (conv1): Conv2d(1024, 256, kernel_size=(1, 1), stride=(1, 1), bias=False)
  (bn1): BatchNorm2d(256, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
  (conv2): Conv2d(256, 256, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1), bias=False)
  (bn2): BatchNorm2d(256, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
  (conv3): Conv2d(256, 1024, kernel_size=(1, 1), stride=(1, 1), bias=False)
  (bn3): BatchNorm2d(1024, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
  (relu): ReLU(inplace)
)
)
(layer4): Sequential(
  (0): Bottleneck(
    (conv1): Conv2d(1024, 512, kernel_size=(1, 1), stride=(1, 1), bias=False)
    (bn1): BatchNorm2d(512, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
    (conv2): Conv2d(512, 512, kernel_size=(3, 3), stride=(2, 2), padding=(1, 1), bias=False)
    (bn2): BatchNorm2d(512, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
    (conv3): Conv2d(512, 2048, kernel_size=(1, 1), stride=(1, 1), bias=False)
    (bn3): BatchNorm2d(2048, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
    (relu): ReLU(inplace)
    (downsample): Sequential(
      (0): Conv2d(1024, 2048, kernel_size=(1, 1), stride=(2, 2), bias=False)
      (1): BatchNorm2d(2048, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
    )
  )
  (1): Bottleneck(
    (conv1): Conv2d(2048, 512, kernel_size=(1, 1), stride=(1, 1), bias=False)
    (bn1): BatchNorm2d(512, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
    (conv2): Conv2d(512, 512, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1), bias=False)
    (bn2): BatchNorm2d(512, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
    (conv3): Conv2d(512, 2048, kernel_size=(1, 1), stride=(1, 1), bias=False)
    (bn3): BatchNorm2d(2048, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
    (relu): ReLU(inplace)
  )
  (2): Bottleneck(
    (conv1): Conv2d(2048, 512, kernel_size=(1, 1), stride=(1, 1), bias=False)
    (bn1): BatchNorm2d(512, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
    (conv2): Conv2d(512, 512, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1), bias=False)
    (bn2): BatchNorm2d(512, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
    (conv3): Conv2d(512, 2048, kernel_size=(1, 1), stride=(1, 1), bias=False)
    (bn3): BatchNorm2d(2048, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
    (relu): ReLU(inplace)
  )
)

```

```

    )
)
(avgpool): AvgPool2d(kernel_size=7, stride=1, padding=0)
(fc): Linear(in_features=2048, out_features=133, bias=True)
)

```

Question 5: Outline the steps you took to get to your final CNN architecture and your reasoning at each step. Describe why you think the architecture is suitable for the current problem.

Answer: The Resnet model was chosen for the sake of classification as a proper candidate for image classification.

1.1.14 (IMPLEMENTATION) Specify Loss Function and Optimizer

Use the next code cell to specify a [loss function](#) and [optimizer](#). Save the chosen loss function as `criterion_transfer`, and the optimizer as `optimizer_transfer` below.

```

In [19]: criterion_transfer = nn.CrossEntropyLoss()
         optimizer_transfer = optimization.SGD(model_transfer.fc.parameters(), lr=0.001)

```

1.1.15 (IMPLEMENTATION) Train and Validate the Model

Train and validate your model in the code cell below. [Save the final model parameters](#) at filepath `'model_transfer.pt'`.

```

In [20]: def train(n_epochs, loaders, model, optimizer, criterion, use_cuda, save_path):
         """returns trained model"""

         valid_loss_min = np.Inf # initialization
         print_after = 100

         print("Training is now starting! Where is your popcorn :}\n")

         for epoch in range(1, n_epochs+1):
             train_loss = 0.0
             valid_loss = 0.0

             #####
             # train the model #
             #####
             model.train()

             for batch_idx, (data, target) in enumerate(loaders['train']):

                 if use_cuda:
                     data, target = data.cuda(), target.cuda()

                 optimizer.zero_grad()

```



```

        output = model(data)
        loss = criterion(output, target)
        loss.backward()
        optimizer.step()
        train_loss += ((1 / (batch_idx + 1)) * (loss.data - train_loss))

    if batch_idx % print_after == 0:
        print(f'\tEpoch #{epoch}, Iteration #{batch_idx+1}, Loss: {train_loss}')

#####
# validate the model #
#####
model.eval()
for batch_idx, (data, target) in enumerate(loaders['valid']):

    if use_cuda:
        data, target = data.cuda(), target.cuda()

    output = model(data)
    loss = criterion(output, target)
    valid_loss += ((1 / (batch_idx + 1)) * (loss.data - valid_loss))

print('Epoch: {} \tTraining Loss: {:.6f} \tValidation Loss: {:.6f}'.format(
    epoch,
    train_loss,
    valid_loss
))

if valid_loss < valid_loss_min:
    torch.save(model.state_dict(), save_path)
    print('Validation loss decreased ({:.6f} --> {:.6f}). Saving model ...'.format(
        valid_loss_min, valid_loss))
    torch.save(model.state_dict(), save_path)
print("\n")

return model

# train the model
trained_epochs = 33
model_file = 'model_transfer.pt'

train(trained_epochs, loaders_transfer, model_transfer, optimizer_transfer, criterion_transfer)

print("Training is now completed! Break is over!")

```

Training is now starting! Where is your popcorn :}

Epoch #1, Iteration #1, Loss: 4.871053218841553
Epoch #1, Iteration #101, Loss: 4.906985759735107
Epoch #1, Iteration #201, Loss: 4.86801815032959
Epoch #1, Iteration #301, Loss: 4.8347578048706055
Epoch: 1 Training Loss: 4.822406 Validation Loss: 4.641018
Validation loss decreased (inf --> 4.641018). Saving model ...

Epoch #2, Iteration #1, Loss: 4.727107048034668
Epoch #2, Iteration #101, Loss: 4.664030075073242
Epoch #2, Iteration #201, Loss: 4.637318134307861
Epoch #2, Iteration #301, Loss: 4.606761932373047
Epoch: 2 Training Loss: 4.596100 Validation Loss: 4.392936
Validation loss decreased (4.641018 --> 4.392936). Saving model ...

Epoch #3, Iteration #1, Loss: 4.631146430969238
Epoch #3, Iteration #101, Loss: 4.442633628845215
Epoch #3, Iteration #201, Loss: 4.42738151550293
Epoch #3, Iteration #301, Loss: 4.404378890991211
Epoch: 3 Training Loss: 4.398647 Validation Loss: 4.143116
Validation loss decreased (4.392936 --> 4.143116). Saving model ...

Epoch #4, Iteration #1, Loss: 4.254113674163818
Epoch #4, Iteration #101, Loss: 4.271188259124756
Epoch #4, Iteration #201, Loss: 4.251242160797119
Epoch #4, Iteration #301, Loss: 4.222027778625488
Epoch: 4 Training Loss: 4.212747 Validation Loss: 3.908545
Validation loss decreased (4.143116 --> 3.908545). Saving model ...

Epoch #5, Iteration #1, Loss: 4.1742753982543945
Epoch #5, Iteration #101, Loss: 4.111791133880615
Epoch #5, Iteration #201, Loss: 4.079235076904297
Epoch #5, Iteration #301, Loss: 4.048720836639404
Epoch: 5 Training Loss: 4.040592 Validation Loss: 3.701074
Validation loss decreased (3.908545 --> 3.701074). Saving model ...

Epoch #6, Iteration #1, Loss: 4.139285564422607
Epoch #6, Iteration #101, Loss: 3.9245238304138184
Epoch #6, Iteration #201, Loss: 3.9014439582824707
Epoch #6, Iteration #301, Loss: 3.8797218799591064
Epoch: 6 Training Loss: 3.872854 Validation Loss: 3.510882
Validation loss decreased (3.701074 --> 3.510882). Saving model ...

Epoch #7, Iteration #1, Loss: 3.7922465801239014
Epoch #7, Iteration #101, Loss: 3.7299368381500244
Epoch #7, Iteration #201, Loss: 3.731947898864746
Epoch #7, Iteration #301, Loss: 3.716610908508301
Epoch: 7 Training Loss: 3.707788 Validation Loss: 3.314846
Validation loss decreased (3.510882 --> 3.314846). Saving model ...

Epoch #8, Iteration #1, Loss: 3.778322219848633
Epoch #8, Iteration #101, Loss: 3.6130785942077637
Epoch #8, Iteration #201, Loss: 3.5760936737060547
Epoch #8, Iteration #301, Loss: 3.564549684524536
Epoch: 8 Training Loss: 3.554457 Validation Loss: 3.132590
Validation loss decreased (3.314846 --> 3.132590). Saving model ...

Epoch #9, Iteration #1, Loss: 3.547138214111328
Epoch #9, Iteration #101, Loss: 3.457639455795288
Epoch #9, Iteration #201, Loss: 3.4368507862091064
Epoch #9, Iteration #301, Loss: 3.417843818664551
Epoch: 9 Training Loss: 3.418759 Validation Loss: 2.957952
Validation loss decreased (3.132590 --> 2.957952). Saving model ...

Epoch #10, Iteration #1, Loss: 3.3525230884552
Epoch #10, Iteration #101, Loss: 3.3215994834899902
Epoch #10, Iteration #201, Loss: 3.308624505996704
Epoch #10, Iteration #301, Loss: 3.2894132137298584
Epoch: 10 Training Loss: 3.284708 Validation Loss: 2.781582
Validation loss decreased (2.957952 --> 2.781582). Saving model ...

Epoch #11, Iteration #1, Loss: 3.0038931369781494
Epoch #11, Iteration #101, Loss: 3.223510503768921
Epoch #11, Iteration #201, Loss: 3.1848623752593994
Epoch #11, Iteration #301, Loss: 3.1724538803100586
Epoch: 11 Training Loss: 3.166575 Validation Loss: 2.639552
Validation loss decreased (2.781582 --> 2.639552). Saving model ...

Epoch #12, Iteration #1, Loss: 2.9938621520996094
Epoch #12, Iteration #101, Loss: 3.0716328620910645
Epoch #12, Iteration #201, Loss: 3.0488555431365967
Epoch #12, Iteration #301, Loss: 3.0349502563476562
Epoch: 12 Training Loss: 3.035010 Validation Loss: 2.518097
Validation loss decreased (2.639552 --> 2.518097). Saving model ...

Epoch #13, Iteration #1, Loss: 3.085958242416382
Epoch #13, Iteration #101, Loss: 2.9398319721221924
Epoch #13, Iteration #201, Loss: 2.9269754886627197
Epoch #13, Iteration #301, Loss: 2.9255728721618652
Epoch: 13 Training Loss: 2.925945 Validation Loss: 2.404809
Validation loss decreased (2.518097 --> 2.404809). Saving model ...

Epoch #14, Iteration #1, Loss: 2.5344018936157227
Epoch #14, Iteration #101, Loss: 2.83162522315979
Epoch #14, Iteration #201, Loss: 2.8394277095794678
Epoch #14, Iteration #301, Loss: 2.833806037902832
Epoch: 14 Training Loss: 2.832696 Validation Loss: 2.282537
Validation loss decreased (2.404809 --> 2.282537). Saving model ...

Epoch #15, Iteration #1, Loss: 2.778444290161133
Epoch #15, Iteration #101, Loss: 2.713353395462036
Epoch #15, Iteration #201, Loss: 2.742091655731201
Epoch #15, Iteration #301, Loss: 2.726854085922241
Epoch: 15 Training Loss: 2.732027 Validation Loss: 2.167162
Validation loss decreased (2.282537 --> 2.167162). Saving model ...

Epoch #16, Iteration #1, Loss: 2.789745569229126
Epoch #16, Iteration #101, Loss: 2.6605262756347656
Epoch #16, Iteration #201, Loss: 2.6723146438598633
Epoch #16, Iteration #301, Loss: 2.66235613822937
Epoch: 16 Training Loss: 2.656143 Validation Loss: 2.088530
Validation loss decreased (2.167162 --> 2.088530). Saving model ...

Epoch #17, Iteration #1, Loss: 2.668334484100342
Epoch #17, Iteration #101, Loss: 2.5816774368286133
Epoch #17, Iteration #201, Loss: 2.566585063934326
Epoch #17, Iteration #301, Loss: 2.5545690059661865
Epoch: 17 Training Loss: 2.554657 Validation Loss: 1.977954
Validation loss decreased (2.088530 --> 1.977954). Saving model ...

Epoch #18, Iteration #1, Loss: 2.268298625946045
Epoch #18, Iteration #101, Loss: 2.5113885402679443
Epoch #18, Iteration #201, Loss: 2.506760358810425
Epoch #18, Iteration #301, Loss: 2.48972749710083
Epoch: 18 Training Loss: 2.489397 Validation Loss: 1.886684
Validation loss decreased (1.977954 --> 1.886684). Saving model ...

Epoch #19, Iteration #1, Loss: 2.3887739181518555
Epoch #19, Iteration #101, Loss: 2.441436529159546
Epoch #19, Iteration #201, Loss: 2.429358720779419
Epoch #19, Iteration #301, Loss: 2.4147322177886963
Epoch: 19 Training Loss: 2.410990 Validation Loss: 1.839738
Validation loss decreased (1.886684 --> 1.839738). Saving model ...

Epoch #20, Iteration #1, Loss: 2.40118145942688
Epoch #20, Iteration #101, Loss: 2.367600440979004
Epoch #20, Iteration #201, Loss: 2.369434356689453
Epoch #20, Iteration #301, Loss: 2.3599424362182617
Epoch: 20 Training Loss: 2.355678 Validation Loss: 1.743631
Validation loss decreased (1.839738 --> 1.743631). Saving model ...

Epoch #21, Iteration #1, Loss: 2.154035806655884
Epoch #21, Iteration #101, Loss: 2.303112030029297
Epoch #21, Iteration #201, Loss: 2.2751400470733643
Epoch #21, Iteration #301, Loss: 2.2692527770996094
Epoch: 21 Training Loss: 2.271598 Validation Loss: 1.691054
Validation loss decreased (1.743631 --> 1.691054). Saving model ...

Epoch #22, Iteration #1, Loss: 2.380399465560913
Epoch #22, Iteration #101, Loss: 2.24202561378479
Epoch #22, Iteration #201, Loss: 2.215045928955078
Epoch #22, Iteration #301, Loss: 2.2346031665802
Epoch: 22 Training Loss: 2.217704 Validation Loss: 1.633183
Validation loss decreased (1.691054 --> 1.633183). Saving model ...

Epoch #23, Iteration #1, Loss: 2.246616840362549
Epoch #23, Iteration #101, Loss: 2.166900873184204
Epoch #23, Iteration #201, Loss: 2.159238576889038
Epoch #23, Iteration #301, Loss: 2.168975353240967
Epoch: 23 Training Loss: 2.164917 Validation Loss: 1.560359
Validation loss decreased (1.633183 --> 1.560359). Saving model ...

Epoch #24, Iteration #1, Loss: 2.1551430225372314
Epoch #24, Iteration #101, Loss: 2.0878403186798096
Epoch #24, Iteration #201, Loss: 2.1098127365112305
Epoch #24, Iteration #301, Loss: 2.0934946537017822
Epoch: 24 Training Loss: 2.091671 Validation Loss: 1.502749
Validation loss decreased (1.560359 --> 1.502749). Saving model ...

Epoch #25, Iteration #1, Loss: 2.0316720008850098
Epoch #25, Iteration #101, Loss: 2.043164014816284
Epoch #25, Iteration #201, Loss: 2.0630781650543213
Epoch #25, Iteration #301, Loss: 2.063009738922119
Epoch: 25 Training Loss: 2.057938 Validation Loss: 1.465889
Validation loss decreased (1.502749 --> 1.465889). Saving model ...

Epoch #26, Iteration #1, Loss: 2.0052175521850586
Epoch #26, Iteration #101, Loss: 2.027733325958252
Epoch #26, Iteration #201, Loss: 2.0338637828826904
Epoch #26, Iteration #301, Loss: 2.0387582778930664
Epoch: 26 Training Loss: 2.033662 Validation Loss: 1.407057
Validation loss decreased (1.465889 --> 1.407057). Saving model ...

Epoch #27, Iteration #1, Loss: 1.824806571006775
Epoch #27, Iteration #101, Loss: 2.0029516220092773
Epoch #27, Iteration #201, Loss: 1.9999046325683594
Epoch #27, Iteration #301, Loss: 1.983224630355835
Epoch: 27 Training Loss: 1.981071 Validation Loss: 1.370805
Validation loss decreased (1.407057 --> 1.370805). Saving model ...

Epoch #28, Iteration #1, Loss: 2.419095516204834
Epoch #28, Iteration #101, Loss: 1.9368305206298828
Epoch #28, Iteration #201, Loss: 1.938778042793274
Epoch #28, Iteration #301, Loss: 1.949415683746338
Epoch: 28 Training Loss: 1.942551 Validation Loss: 1.322099
Validation loss decreased (1.370805 --> 1.322099). Saving model ...

Epoch #29, Iteration #1, Loss: 2.0739150047302246
Epoch #29, Iteration #101, Loss: 1.8892093896865845
Epoch #29, Iteration #201, Loss: 1.879821538925171
Epoch #29, Iteration #301, Loss: 1.8724563121795654
Epoch: 29 Training Loss: 1.877446 Validation Loss: 1.279112
Validation loss decreased (1.322099 --> 1.279112). Saving model ...

Epoch #30, Iteration #1, Loss: 2.053122043609619
Epoch #30, Iteration #101, Loss: 1.8621724843978882
Epoch #30, Iteration #201, Loss: 1.8656651973724365
Epoch #30, Iteration #301, Loss: 1.8602133989334106
Epoch: 30 Training Loss: 1.863457 Validation Loss: 1.239220
Validation loss decreased (1.279112 --> 1.239220). Saving model ...

```

Epoch #31, Iteration #1, Loss: 2.3706531524658203
Epoch #31, Iteration #101, Loss: 1.8474818468093872
Epoch #31, Iteration #201, Loss: 1.8314679861068726
Epoch #31, Iteration #301, Loss: 1.8264859914779663
Epoch: 31      Training Loss: 1.822910      Validation Loss: 1.226833
Validation loss decreased (1.239220 --> 1.226833). Saving model ...

```

```

Epoch #32, Iteration #1, Loss: 1.4818055629730225
Epoch #32, Iteration #101, Loss: 1.7899476289749146
Epoch #32, Iteration #201, Loss: 1.7933275699615479
Epoch #32, Iteration #301, Loss: 1.8014802932739258
Epoch: 32      Training Loss: 1.794790      Validation Loss: 1.195498
Validation loss decreased (1.226833 --> 1.195498). Saving model ...

```

```

Epoch #33, Iteration #1, Loss: 2.1339802742004395
Epoch #33, Iteration #101, Loss: 1.7751797437667847
Epoch #33, Iteration #201, Loss: 1.778869390487671
Epoch #33, Iteration #301, Loss: 1.7815287113189697
Epoch: 33      Training Loss: 1.780451      Validation Loss: 1.145228
Validation loss decreased (1.195498 --> 1.145228). Saving model ...

```

Training is now completed! Break is over!

1.1.16 (IMPLEMENTATION) Test the Model

Try out your model on the test dataset of dog images. Use the code cell below to calculate and print the test loss and accuracy. Ensure that your test accuracy is greater than 60%.

```
In [21]: test(loaders_transfer, model_transfer, criterion_transfer, use_cuda)
```

```
Test Loss: 1.211861
```

```
Test Accuracy: 78% (653/836)
```

1.1.17 (IMPLEMENTATION) Predict Dog Breed with the Model

Write a function that takes an image path as input and returns the dog breed (Affenpinscher, Afghan hound, etc) that is predicted by your model.

```
In [35]: #from PIL import Image
         #import torchvision.transforms as transforms
```

```

import cv2
import matplotlib.pyplot as plt
%matplotlib inline

class_names = [item[4:].replace("_", " ") for item in loaders_transfer['train'].dataset
loaders_transfer['train'].dataset.classes[:10]
class_names[:10]

def get_image(img_path):
    img = Image.open(img_path).convert('RGB')
    prediction_transform = transforms.Compose([transforms.Resize(size=(224, 224)),trans
    img = prediction_transform(img)[:3,:,:].unsqueeze(0)
    return img

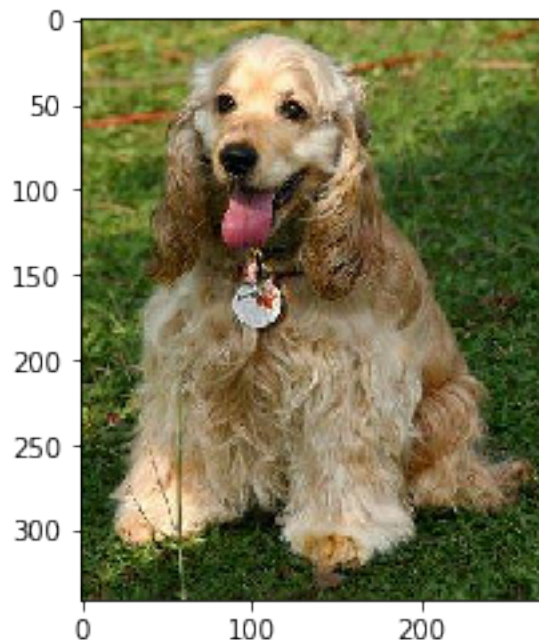
def predict_breed_transfer(model, class_names, img_path):
    img = get_image(img_path)
    pred = torch.argmax(model(img))
    return class_names[pred]

```


```

In [36]: # test
for img_file in os.listdir('/data/dog_images/test/053.Cocker_spaniel/'):
    img = cv2.imread(img_path)
    cv_rgb = cv2.cvtColor(img, cv2.COLOR_BGR2RGB)
    plt.imshow(cv_rgb)
    plt.show()
    img_path = os.path.join('/data/dog_images/test/053.Cocker_spaniel/', img_file)
    predition = predict_breed_transfer(model_transfer, class_names, img_path)
    print("predition breed: {}".format(predition))


```




prediction breed: Cocker spaniel




prediction breed: Cocker spaniel



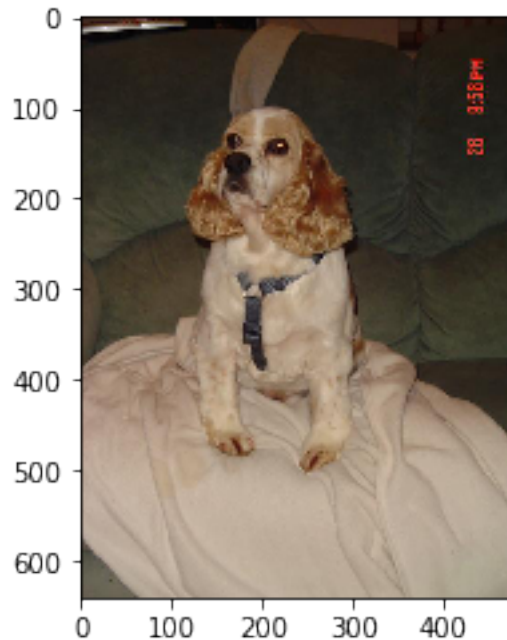
prediction breed: Cavalier king charles spaniel



prediction breed: English cocker spaniel



```
prediction breed: English cocker spaniel
```



```
prediction breed: English cocker spaniel
```

Step 5: Write your Algorithm

Write an algorithm that accepts a file path to an image and first determines whether the image contains a human, dog, or neither. Then, - if a **dog** is detected in the image, return the predicted breed. - if a **human** is detected in the image, return the resembling dog breed. - if **neither** is detected in the image, provide output that indicates an error.

You are welcome to write your own functions for detecting humans and dogs in images, but feel free to use the `face_detector` and `human_detector` functions developed above. You are **required** to use your CNN from Step 4 to predict dog breed.

Some sample output for our algorithm is provided below, but feel free to design your own user experience!

1.1.18 (IMPLEMENTATION) Write your Algorithm

```
In [46]: def dog_app(img_path):  
         img = Image.open(img_path)
```



Sample Human Output

```
plt.imshow(img)
plt.show()
if dog_detector(img_path) is True:
    #prediction = predict_breed_transfer(model_transfer, class_names, img_path)
    print("This dog looks like a {0}".format(predict_breed_transfer(model_transfer,
elif face_detector(img_path) > 0:
    prediction = predict_breed_transfer(model_transfer, class_names, img_path)
    print("This human looks like the dog breed {0}".format(prediction))
else:
    print("Neither dog nor human was detected by the model")
```

Step 6: Test Your Algorithm

In this section, you will take your new algorithm for a spin! What kind of dog does the algorithm think that *you* look like? If you have a dog, does it predict your dog's breed accurately? If you have a cat, does it mistakenly think that your cat is a dog?


1.1.19 (IMPLEMENTATION) Test Your Algorithm on Sample Images!

Test your algorithm at least six images on your computer. Feel free to use any images you like. Use at least two human and two dog images.


Question 6: Is the output better than you expected :) ? Or worse :(? Provide at least three possible points of improvement for your algorithm.

Answer: (Three possible points for improvement): more image for train the model, increase the overall performance of the model, and parameterize the model more.


```
In [60]: num_human_img = 3
         num_dog_img = 4
         index_slice = 3
         for file in np.hstack((human_files[index_slice:(num_human_img+index_slice)], dog_files[
```




This human looks like the dog breed Poodle




This human looks like the dog breed German pinscher



This human looks like the dog breed Dachshund



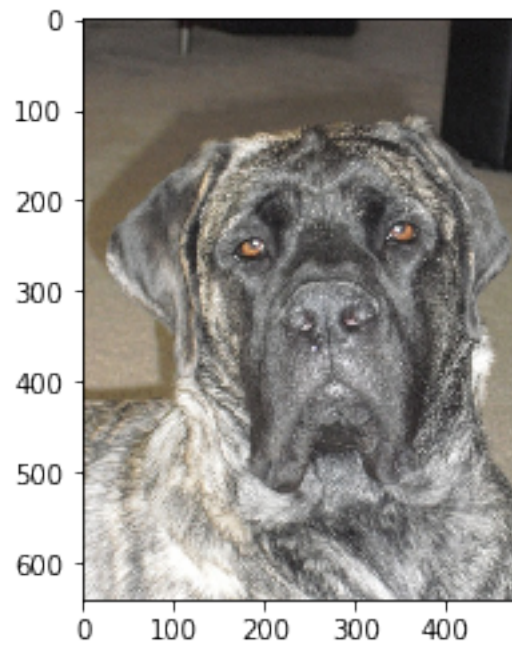
This dog looks like a Mastiff



This dog looks like a Mastiff



This dog looks like a Mastiff



This dog looks like a Mastiff

In []: